

# dplyr and pipes: the basics

Sean C. Anderson, [sean@seananderson.ca](mailto:sean@seananderson.ca)

September 16, 2014

## Contents

Introduction	1
The data	2
Looking at the data	2
Selecting columns	3
Filtering rows	6
Arranging rows	7
Mutating columns	8
Summarising columns	9
Piping data	9
A trivial pipe example	10
A silly dplyr example with pipes	10
An awesome example	10

## Introduction

The [dplyr](#) R package is awesome. Pipes from the [magrittr](#) R package are awesome. Put the two together and you have one of the most exciting things to happen to R in a long time.

dplyr is [Hadley Wickham's](#) re-imagined plyr package (with underlying C++ secret sauce co-written by [Romain Francois](#)). plyr 2.0 if you will. It does less than dplyr, but what it does it does more elegantly and much more quickly.

dplyr is built around 5 verbs. These verbs make up the majority of the data manipulation you tend to do. You might need to:

*Select* certain columns of data.

*Filter* your data to select specific rows.

*Arrange* the rows of your data into an order.

*Mutate* your data frame to contain new columns.

*Summarise* chunks of your data in some way.

Let's look at how those work.

## The data

We're going to work with a dataset of mammal life-history, geography, and ecology traits from the PanTHERIA database:

Jones, K.E., *et al.* PanTHERIA: a species-level database of life history, ecology, and geography of extant and recently extinct mammals. Ecology 90:2648. <http://esapubs.org/archive/ecol/E090/184/>

First we'll download the data:

```
pantheria <-  
  "http://esapubs.org/archive/ecol/E090/184/PanTHERIA_1-0_WR05_Aug2008.txt"  
download.file(pantheria, destfile = "mammals.txt")
```

Next we'll read it in and simplify it. This gets a bit ugly, but you can safely just run this code chunk and ignore the details:

```
mammals <- read.table("mammals.txt", sep = "\t", header = TRUE,  
  stringsAsFactors = FALSE)  
names(mammals) <- sub("X[0-9.]+", "", names(mammals))  
names(mammals) <- sub("MSW05_", "", names(mammals))  
mammals <- dplyr::select(mammals, Order, Binomial, AdultBodyMass_g,  
  AdultHeadBodyLen_mm, HomeRange_km2, LitterSize)  
names(mammals) <- gsub("[A-Z]", "_\\L\\1", names(mammals), perl = TRUE)  
names(mammals) <- gsub("^_", "", names(mammals), perl = TRUE)  
mammals[mammals == -999] <- NA  
names(mammals)[names(mammals) == "binomial"] <- "species"  
mammals <- dplyr::tbl_df(mammals) # for prettier printing
```

Next we'll load the dplyr package:

```
library(dplyr)
```

## Looking at the data

Data frames look a bit different in dplyr. Above, I called the `tbl_df()` function on our data. This provides more useful printing of data frames in the console. Ever accidentally printed a massive data frame in the console before? Yeah... this avoids that. You don't need to change your data to a data frame `tbl` first — the dplyr functions will automatically convert your data when you call them. This is what the data look like on the console:

```
mammals
```

```
## Source: local data frame [5,416 x 6]  
##
```

```
##           order           species adult_body_mass_g
## 1 Artiodactyla Camelus dromedarius      492714
## 2 Carnivora      Canis adustus        10392
## 3 Carnivora      Canis aureus         9659
## 4 Carnivora      Canis latrans        11989
## 5 Carnivora      Canis lupus          31757
## 6 Artiodactyla      Bos frontalis      800143
## 7 Artiodactyla      Bos grunniens      500000
## 8 Artiodactyla      Bos javanicus      635974
## 9 Primates Callicebus cupreus         1117
## 10 Primates Callicebus discolor          NA
## ..           ...           ...
## Variables not shown: adult_head_body_len_mm (dbl), home_range_km2 (dbl),
##   litter_size (dbl)
```

dplyr also provides a function `glimpse()` that makes it easy to look at our data in a transposed view. It's similar to the `str()` (structure) function, but has a few advantages (see `?glimpse`).

```
glimpse(mammals)
```

```
## Variables:
## $ order      (chr) "Artiodactyla", "Carnivora", "Carnivora..."
## $ species     (chr) "Camelus dromedarius", "Canis adustus",...
## $ adult_body_mass_g (dbl) 492714.5, 10392.5, 9658.7, 11989.1, 317...
## $ adult_head_body_len_mm (dbl) NA, 745.3, 827.5, 872.4, 1055.0, 2700.0...
## $ home_range_km2 (dbl) 196.32000, 1.01000, 2.95000, 18.88000, ...
## $ litter_size (dbl) 0.98, 4.50, 3.74, 5.72, 4.98, 1.22, 1.0...
```

## Selecting columns

`select()` lets you subset by columns. This is similar to `subset()` in base R, but it also allows for some fancy use of helper functions such as `contains()`, `starts_with()` and `ends_with()`. I think these examples are self explanatory, so I'll just include them here:

```
select(mammals, adult_head_body_len_mm)
```

```
## Source: local data frame [5,416 x 1]
##
##   adult_head_body_len_mm
## 1 NA
## 2 745.3
## 3 827.5
## 4 872.4
## 5 1055.0
## 6 2700.0
## 7 NA
## 8 2075.0
## 9 355.0
## 10 NA
## ..           ...
```

```
select(mammals, adult_head_body_len_mm, litter_size)
```

```
## Source: local data frame [5,416 x 2]
##
##   adult_head_body_len_mm litter_size
## 1                      NA         0.98
## 2                   745.3         4.50
## 3                   827.5         3.74
## 4                   872.4         5.72
## 5                  1055.0         4.98
## 6                  2700.0         1.22
## 7                      NA         1.00
## 8                  2075.0         1.22
## 9                   355.0         1.01
## 10                     NA          NA
## ..                      ...         ...
```

```
select(mammals, adult_head_body_len_mm:litter_size)
```

```
## Source: local data frame [5,416 x 3]
##
##   adult_head_body_len_mm home_range_km2 litter_size
## 1                      NA         196.32         0.98
## 2                   745.3          1.01         4.50
## 3                   827.5          2.95         3.74
## 4                   872.4         18.88         5.72
## 5                  1055.0        159.86         4.98
## 6                  2700.0          NA         1.22
## 7                      NA          NA         1.00
## 8                  2075.0          NA         1.22
## 9                   355.0          NA         1.01
## 10                     NA          NA          NA
## ..                      ...         ...         ...
```

```
select(mammals, -adult_head_body_len_mm)
```

```
## Source: local data frame [5,416 x 5]
##
##   order      species adult_body_mass_g home_range_km2
## 1 Artiodactyla Camelus dromedarius    492714        196.32
## 2 Carnivora    Canis adustus      10392          1.01
## 3 Carnivora    Canis aureus       9659          2.95
## 4 Carnivora    Canis latrans     11989         18.88
## 5 Carnivora    Canis lupus      31757        159.86
## 6 Artiodactyla Bos frontalis    800143          NA
## 7 Artiodactyla Bos grunniens  500000          NA
## 8 Artiodactyla Bos javanicus   635974          NA
## 9 Primates    Callicebus cupreus    1117          NA
## 10 Primates   Callicebus discolor      NA          NA
## ..      ...      ...      ...      ...
## Variables not shown: litter_size (dbl)
```

```
select(mammals, contains("body"))
```

```
## Source: local data frame [5,416 x 2]
##
##   adult_body_mass_g adult_head_body_len_mm
## 1           492714                NA
## 2           10392                745.3
## 3            9659                827.5
## 4           11989                872.4
## 5           31757               1055.0
## 6          800143               2700.0
## 7          500000                NA
## 8          635974               2075.0
## 9            1117               355.0
## 10             NA                NA
## ..           ...                ...
```

```
select(mammals, starts_with("adult"))
```

```
## Source: local data frame [5,416 x 2]
##
##   adult_body_mass_g adult_head_body_len_mm
## 1           492714                NA
## 2           10392                745.3
## 3            9659                827.5
## 4           11989                872.4
## 5           31757               1055.0
## 6          800143               2700.0
## 7          500000                NA
## 8          635974               2075.0
## 9            1117               355.0
## 10             NA                NA
## ..           ...                ...
```

```
select(mammals, ends_with("g"))
```

```
## Source: local data frame [5,416 x 1]
##
##   adult_body_mass_g
## 1           492714
## 2           10392
## 3            9659
## 4           11989
## 5           31757
## 6          800143
## 7          500000
## 8          635974
## 9            1117
## 10             NA
## ..           ...
```

```
select(mammals, 1:3)
```

```
## Source: local data frame [5,416 x 3]
##
##      order      species adult_body_mass_g
## 1 Artiodactyla Camelus dromedarius      492714
## 2 Carnivora      Canis adustus        10392
## 3 Carnivora      Canis aureus         9659
## 4 Carnivora      Canis latrans       11989
## 5 Carnivora      Canis lupus        31757
## 6 Artiodactyla      Bos frontalis     800143
## 7 Artiodactyla      Bos grunniens    500000
## 8 Artiodactyla      Bos javanicus    635974
## 9 Primates Callicebus cupreus        1117
## 10 Primates Callicebus discolor         NA
## ..      ...      ...      ...
```

## Filtering rows

`filter()` lets you subset by rows. You can use any valid logical statements:

```
filter(mammals, adult_body_mass_g > 1e7)[ , 1:3]
```

```
## Source: local data frame [12 x 3]
##
##      order      species adult_body_mass_g
## 1 Cetacea      Caperea marginata    32000000
## 2 Cetacea Balaenoptera musculus    154321304
## 3 Cetacea Balaenoptera physalus    47506008
## 4 Cetacea Balaena mysticetus    79691179
## 5 Cetacea Balaenoptera borealis    22106252
## 6 Cetacea Balaenoptera edeni    20000000
## 7 Cetacea Berardius bairdii    11380000
## 8 Cetacea Eschrichtius robustus    27324024
## 9 Cetacea Eubalaena australis    23000000
## 10 Cetacea Eubalaena glacialis    23000000
## 11 Cetacea Megaptera novaeangliae    30000000
## 12 Cetacea Physeter catodon    14540960
```

```
filter(mammals, species == "Balaena mysticetus")
```

```
## Source: local data frame [1 x 6]
##
##      order      species adult_body_mass_g adult_head_body_len_mm
## 1 Cetacea Balaena mysticetus    79691179    12187
## Variables not shown: home_range_km2 (dbl), litter_size (dbl)
```

```
filter(mammals, order == "Carnivora" & adult_body_mass_g < 200)
```

```
## Source: local data frame [3 x 6]
##
##      order      species adult_body_mass_g adult_head_body_len_mm
## 1 Carnivora Mustela altaica           180.24           243.5
## 2 Carnivora Mustela frenata           190.03           229.3
## 3 Carnivora Mustela nivalis            78.45           188.2
## Variables not shown: home_range_km2 (dbl), litter_size (dbl)
```

## Arranging rows

`arrange()` lets you order the rows by one or more columns in ascending or descending order. I'm selecting the first three columns only to make the output easier to read:

```
arrange(mammals, adult_body_mass_g)[ , 1:3]
```

```
## Source: local data frame [5,416 x 3]
##
##      order      species adult_body_mass_g
## 1  Chiroptera Craseonycteris thonglongyai      1.96
## 2  Chiroptera Kerivoula minuta              2.03
## 3  Soricomorpha Suncus etruscus             2.26
## 4  Soricomorpha Sorex minutissimus          2.46
## 5  Soricomorpha Suncus madagascariensis      2.47
## 6  Soricomorpha Crocidura lusitania         2.48
## 7  Soricomorpha Crocidura planiceps         2.50
## 8  Chiroptera Pipistrellus nanulus          2.51
## 9  Soricomorpha Sorex nanus                 2.57
## 10 Soricomorpha Sorex arizonae             2.70
## ..      ...                ...                ...
```

```
arrange(mammals, desc(adult_body_mass_g))[ , 1:3]
```

```
## Source: local data frame [5,416 x 3]
##
##      order      species adult_body_mass_g
## 1  Cetacea Balaenoptera musculus      154321304
## 2  Cetacea Balaena mysticetus        79691179
## 3  Cetacea Balaenoptera physalus      47506008
## 4  Cetacea Caperea marginata         32000000
## 5  Cetacea Megaptera novaeangliae     30000000
## 6  Cetacea Eschrichtius robustus      27324024
## 7  Cetacea Eubalaena australis        23000000
## 8  Cetacea Eubalaena glacialis        23000000
## 9  Cetacea Balaenoptera borealis      22106252
## 10 Cetacea Balaenoptera edeni         20000000
## ..      ...                ...                ...
```

```
arrange(mammals, order, adult_body_mass_g)[ , 1:3]
```

```
## Source: local data frame [5,416 x 3]
```

```
##
##           order                species adult_body_mass_g
## 1 Afrosoricida      Microgale pusilla          3.40
## 2 Afrosoricida      Microgale parvula          3.53
## 3 Afrosoricida      Geogale aurita            6.69
## 4 Afrosoricida      Microgale fotsifotsy       7.70
## 5 Afrosoricida      Microgale longicaudata      8.08
## 6 Afrosoricida      Microgale brevicaudata      8.99
## 7 Afrosoricida      Microgale principula       10.20
## 8 Afrosoricida      Microgale drouhardi        10.50
## 9 Afrosoricida      Microgale cowani           12.27
## 10 Afrosoricida      Microgale taiva            12.40
## ..              ...                ...          ...
```

## Mutating columns

`mutate()` lets you add new columns. Notice that the new columns you create can build on each other. I will wrap these in `glimpse()` to make the new columns easy to see:

```
glimpse(mutate(mammals, adult_body_mass_kg = adult_body_mass_g / 1000))
```

```
## Variables:
## $ order           (chr) "Artiodactyla", "Carnivora", "Carnivora..."
## $ species         (chr) "Camelus dromedarius", "Canis adustus",...
## $ adult_body_mass_g (dbl) 492714.5, 10392.5, 9658.7, 11989.1, 317...
## $ adult_head_body_len_mm (dbl) NA, 745.3, 827.5, 872.4, 1055.0, 2700.0...
## $ home_range_km2   (dbl) 196.32000, 1.01000, 2.95000, 18.88000, ...
## $ litter_size      (dbl) 0.98, 4.50, 3.74, 5.72, 4.98, 1.22, 1.0...
## $ adult_body_mass_kg (dbl) 492.7145, 10.3925, 9.6587, 11.9891, 31....
```

```
glimpse(mutate(mammals,
  g_per_mm = adult_body_mass_g / adult_head_body_len_mm))
```

```
## Variables:
## $ order           (chr) "Artiodactyla", "Carnivora", "Carnivora..."
## $ species         (chr) "Camelus dromedarius", "Canis adustus",...
## $ adult_body_mass_g (dbl) 492714.5, 10392.5, 9658.7, 11989.1, 317...
## $ adult_head_body_len_mm (dbl) NA, 745.3, 827.5, 872.4, 1055.0, 2700.0...
## $ home_range_km2   (dbl) 196.32000, 1.01000, 2.95000, 18.88000, ...
## $ litter_size      (dbl) 0.98, 4.50, 3.74, 5.72, 4.98, 1.22, 1.0...
## $ g_per_mm         (dbl) NA, 13.9437, 11.6717, 13.7428, 30.1010,...
```

```
glimpse(mutate(mammals,
  g_per_mm = adult_body_mass_g / adult_head_body_len_mm,
  kg_per_mm = g_per_mm / 1000))
```

```
## Variables:
## $ order           (chr) "Artiodactyla", "Carnivora", "Carnivora..."
## $ species         (chr) "Camelus dromedarius", "Canis adustus",...
## $ adult_body_mass_g (dbl) 492714.5, 10392.5, 9658.7, 11989.1, 317...
```



```
## $ adult_head_body_len_mm (dbl) NA, 745.3, 827.5, 872.4, 1055.0, 2700.0...
## $ home_range_km2        (dbl) 196.32000, 1.01000, 2.95000, 18.88000, ...
## $ litter_size           (dbl) 0.98, 4.50, 3.74, 5.72, 4.98, 1.22, 1.0...
## $ g_per_mm              (dbl) NA, 13.9437, 11.6717, 13.7428, 30.1010,...
## $ kg_per_mm             (dbl) NA, 0.0139437, 0.0116717, 0.0137428, 0....
```

## Summarising columns

Finally, `summarise()` lets you calculate summary statistics. On its own `summarise()` isn't that useful, but when combined with `group_by()` you can summarise by chunks of data. This is similar to what you might be familiar with through `ddply()` and `summarise()` from the `plyr` package:

```
summarise(mammals, mean_mass = mean(adult_body_mass_g, na.rm = TRUE))
```

```
## Source: local data frame [1 x 1]
##
##   mean_mass
## 1      177810
```

```
# summarise with group_by:
head(summarise(group_by(mammals, order),
  mean_mass = mean(adult_body_mass_g, na.rm = TRUE)))
```

```
## Source: local data frame [6 x 2]
##
##   order mean_mass
## 1 Afrosoricida 9.476e+01
## 2 Artiodactyla 1.213e+05
## 3 Carnivora 4.739e+04
## 4 Cetacea 7.373e+06
## 5 Chiroptera 5.772e+01
## 6 Cingulata 4.699e+03
```

## Piping data

Pipes take the output from one function and feed it to the first argument of the next function. You may have encountered the Unix pipe `|` before.

The `magrittr` R package contains the pipe function `%>%`. Yes it might look bizarre at first but it makes more sense when you think about it. The R language allows symbols wrapped in `%` to be defined as functions, the `>` helps imply a chain, and you can hit these 2 characters one after the other very quickly on a keyboard by holding down the Shift key. Try it!

Try pronouncing `%>%` “then” whenever you see it. If you want to see the help page, you'll need to wrap it in back ticks like so:

```
?magrittr::`%>%`
```

## A trivial pipe example

Pipes can work with nearly any functions. Let's start with a non-dplyr example:

```
x <- rnorm(10)
x %>% max
```

```
## [1] 1.77
```

```
# is the same thing as:
max(x)
```

```
## [1] 1.77
```

So, we took the value of `x` (what would have been printed on the console), captured it, and fed it to the first argument of `max()`. It's probably not clear why this is cool yet, but hang on.

## A silly dplyr example with pipes

Let's try a single-pipe dplyr example. We'll pipe the `mammals` data frame to the `arrange` function's first argument, and choose to arrange by the `adult_body_mass_g` column:

```
mammals %>% arrange(adult_body_mass_g)
```

```
## Source: local data frame [5,416 x 6]
##
##      order      species adult_body_mass_g
## 1  Chiroptera Craseonycteris thonglongyai    1.96
## 2  Chiroptera Kerivoula minuta            2.03
## 3  Soricomorpha Suncus etruscus            2.26
## 4  Soricomorpha Sorex minutissimus         2.46
## 5  Soricomorpha Suncus madagascariensis     2.47
## 6  Soricomorpha Crocidura lusitania         2.48
## 7  Soricomorpha Crocidura planiceps         2.50
## 8  Chiroptera  Pipistrellus nanulus         2.51
## 9  Soricomorpha Sorex nanus                2.57
## 10 Soricomorpha Sorex arizonae             2.70
## ..      ...
## Variables not shown: adult_head_body_len_mm (dbl), home_range_km2 (dbl),
##   litter_size (dbl)
```

## An awesome example

OK, here's where it gets cool. We can chain dplyr functions in succession. This lets us write data manipulation steps in the order we think of them and avoid creating temporary variables in the middle to capture the output. This works because the output from every dplyr function is a data frame and the first argument of every dplyr function is a data frame.

Say we wanted to find the species with the highest body-mass-to-length ratio:

```
mammals %>%
  mutate(mass_to_length = adult_body_mass_g / adult_head_body_len_mm) %>%
  arrange(desc(mass_to_length)) %>%
  select(species, mass_to_length)
```

```
## Source: local data frame [5,416 x 2]
##
##           species mass_to_length
## 1 Balaena mysticetus      6539
## 2 Balaenoptera musculus    5063
## 3 Megaptera novaeangliae   2334
## 4 Eschrichtius robustus    2309
## 5 Balaenoptera physalus    2302
## 6 Elephas maximus         1704
## 7 Eubalaena glacialis      1654
## 8 Eubalaena australis      1625
## 9 Balaenoptera edeni       1444
## 10 Balaenoptera borealis    1203
## .. ... ..
```

So, we took `mammals`, fed it to `mutate()` to create a mass-length ratio column, arranged the resulting data frame in descending order by that ratio, and selected the columns we wanted to see. This is just the beginning. If you can imagine it, you can string it together. If you want to debug your code, just pull a pipe off the end and run the code down to that step. Or build your analysis up and add successive pipes.

The above is equivalent to:

```
select(
  arrange(
    mutate(mammals,
      mass_to_length = adult_body_mass_g / adult_head_body_len_mm),
    desc(mass_to_length)),
  species, mass_to_length)
```

```
## Source: local data frame [5,416 x 2]
##
##           species mass_to_length
## 1 Balaena mysticetus      6539
## 2 Balaenoptera musculus    5063
## 3 Megaptera novaeangliae   2334
## 4 Eschrichtius robustus    2309
## 5 Balaenoptera physalus    2302
## 6 Elephas maximus         1704
## 7 Eubalaena glacialis      1654
## 8 Eubalaena australis      1625
## 9 Balaenoptera edeni       1444
## 10 Balaenoptera borealis    1203
## .. ... ..
```

But the problem here is that you have to read it inside out, it's easy to miss a bracket, and the arguments get separated from the function (e.g. see `mutate()` and `desc(mass_to_length)`). Plus, this is a rather trivial example. Chain together even more steps and it quickly gets out of hand.

Here's one more example. Let's ask what taxonomic orders have a median litter size greater than 3.

```
mammals %>% group_by(order) %>%
  summarise(median_litter = median(litter_size, na.rm = TRUE)) %>%
  filter(median_litter > 3) %>%
  arrange(desc(median_litter)) %>%
  select(order, median_litter)
```

```
## Source: local data frame [5 x 2]
##
##           order median_litter
## 1 Didelphimorphia      6.595
## 2 Dasyuromorphia      6.190
## 3 Erinaceomorpha      3.870
## 4 Soricomorpha      3.660
## 5 Rodentia      3.280
```

These examples don't even highlight one of the best things about dplyr. It's *really* fast. The internal C++ code makes quick work of massive data frames that would make plyr slow to a crawl.

dplyr can do much more, but the above are the basics of the 5 verbs and pipes. Try them for a bit. Once they click I think they'll revolutionize your data analysis.